

TV Show Shark Tank with Machine Learning

March 6, 2024

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings(action='ignore', category=FutureWarning)
```

1 Reading in the data

```
[2]: explore_df = pd.read_csv('Shark Tank US dataset.csv')
explore_df.head()
```

```
[2]:
```

	Season Number	Season Start	Season End	Episode Number	Pitch Number \
0	1	9-Aug-09	5-Feb-10	1	1
1	1	9-Aug-09	5-Feb-10	1	2
2	1	9-Aug-09	5-Feb-10	1	3
3	1	9-Aug-09	5-Feb-10	1	4
4	1	9-Aug-09	5-Feb-10	1	5

	Original Air Date	Startup Name	Industry \
0	9-Aug-09	AvaTheElephant	Health/Wellness
1	9-Aug-09	Mr.Tod'sPieFactory	Food and Beverage
2	9-Aug-09	Wispots	Business Services
3	9-Aug-09	CollegeFoxesPackingBoxes	Lifestyle/Home
4	9-Aug-09	IonicEar	Software/Tech

	Business Description	Pitchers Gender ... \
0	Ava The Elephant - Baby and Child Care	Female ...
1	Mr. Tod's Pie Factory - Specialty Food	Male ...
2	Wispots - Consumer Services	Male ...
3	College Foxes Packing Boxes - Consumer Services	Male ...
4	Ionic Ear - Novelties	Male ...

	Kevin O Leary Investment Equity	Guest Investment Amount \
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN

4	NaN	NaN
---	-----	-----

	Guest Investment Equity	Guest Name	Barbara Corcoran Present	\
0	NaN	NaN	1.0	
1	NaN	NaN	1.0	
2	NaN	NaN	1.0	
3	NaN	NaN	1.0	
4	NaN	NaN	1.0	

	Mark Cuban Present	Lori Greiner Present	Robert Herjavec Present	\
0	0.0	0.0	1.0	
1	0.0	0.0	1.0	
2	0.0	0.0	1.0	
3	0.0	0.0	1.0	
4	0.0	0.0	1.0	

	Daymond John Present	Kevin O Leary Present
0	1.0	1.0
1	1.0	1.0
2	1.0	1.0
3	1.0	1.0
4	1.0	1.0

[5 rows x 50 columns]

```
[3]: explore_df.shape
```

```
[3]: (1274, 50)
```

```
[4]: explore_df.isna().sum()
```

```
[4]: Season Number      0
Season Start           0
Season End             0
Episode Number        0
Pitch Number          0
Original Air Date      0
Startup Name           0
Industry              0
Business Description   0
Pitchers Gender        7
Pitchers City          772
Pitchers State         528
Pitchers Average Age   936
Entrepreneur Names     495
Company Website        758
Multiple Entrepreneurs 427
```

US Viewership	0
Original Ask Amount	0
Original Offered Equity	0
Valuation Requested	0
Got Deal	0
Total Deal Amount	509
Total Deal Equity	509
Deal Valuation	509
Number of sharks in deal	509
Investment Amount Per Shark	509
Equity Per Shark	509
Royalty Deal	1199
Loan	1222
Barbara Corcoran Investment Amount	1154
Barbara Corcoran Investment Equity	1154
Mark Cuban Investment Amount	1044
Mark Cuban Investment Equity	1044
Lori Greiner Investment Amount	1075
Lori Greiner Investment Equity	1075
Robert Herjavec Investment Amount	1153
Robert Herjavec Investment Equity	1153
Daymond John Investment Amount	1163
Daymond John Investment Equity	1163
Kevin O Leary Investment Amount	1157
Kevin O Leary Investment Equity	1157
Guest Investment Amount	1169
Guest Investment Equity	1169
Guest Name	1169
Barbara Corcoran Present	376
Mark Cuban Present	373
Lori Greiner Present	373
Robert Herjavec Present	377
Daymond John Present	376
Kevin O Leary Present	376

dtype: int64

```
[5]: kept_df = explore_df[['Industry', 'Pitchers Gender', 'Original Ask Amount',
    ↪ 'Total Deal Amount',
    ↪ 'Got Deal', 'Barbara Corcoran Present', 'Mark Cuban
    ↪ Present', 'Lori Greiner Present',
    ↪ 'Robert Herjavec Present', 'Daymond John Present', 'Kevin
    ↪ O Leary Present', 'Pitchers State']]

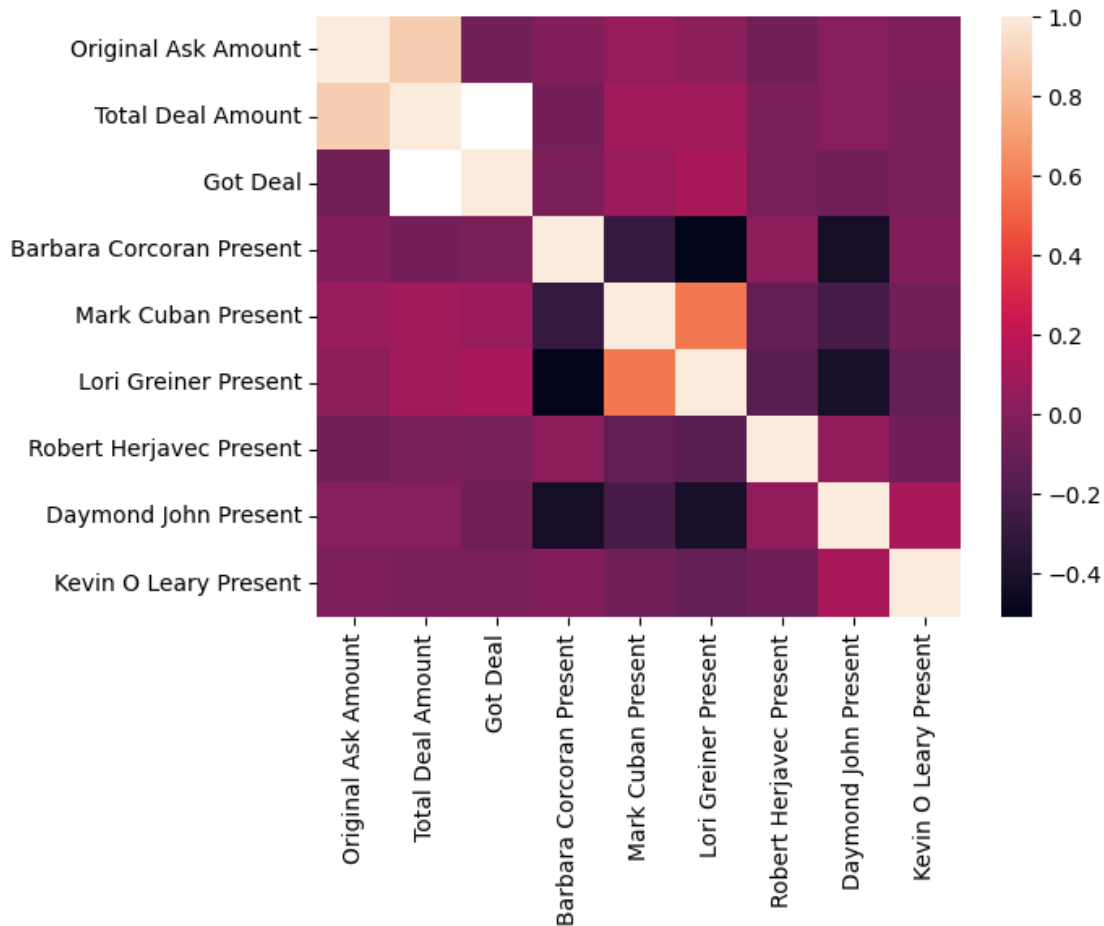
kept_df.shape
```

```
[5]: (1274, 12)
```

```
[6]: initial_drop = kept_df.dropna()
initial_drop.shape
```

```
[6]: (405, 12)
```

```
[7]: # Calculating the correlations and then visualizing the relationships
corr_df = kept_df.corr()
sns.heatmap(corr_df)
plt.show()
```



“Total Deal Amount” has a perfect relationship with “Got Deal” and a very high correlation with “Original Ask Amount”, “Total Deal Amount” will not be used of the model.

```
[8]: # Dropping high correlation feature
v2_kept_df = kept_df.drop('Total Deal Amount', axis=1)
v2_kept_df.shape
```


```
[8]: (1274, 11)
```

```
[9]: dropping_nulls_df = v2_kept_df.dropna()
dropping_nulls_df.shape
```

```
[9]: (734, 11)
```

2 Visualizing the data

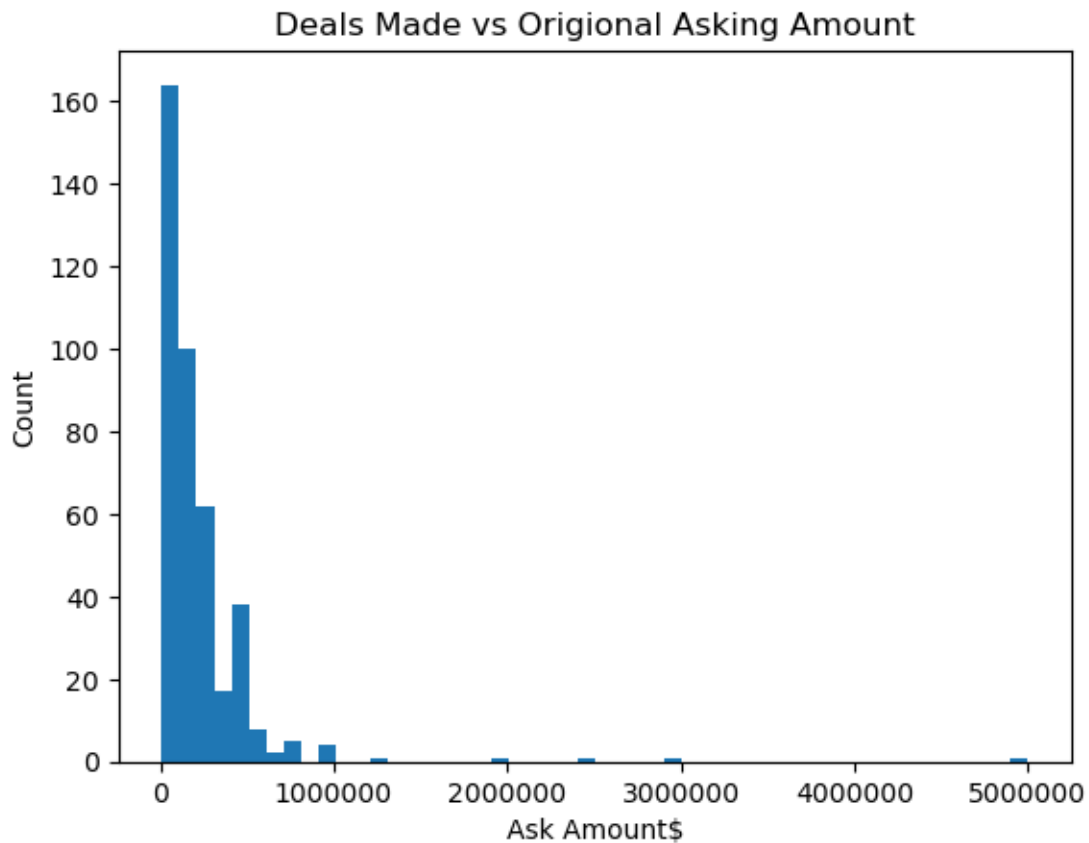
```
[10]: sns.pairplot(dropping_nulls_df)
plt.show()
```



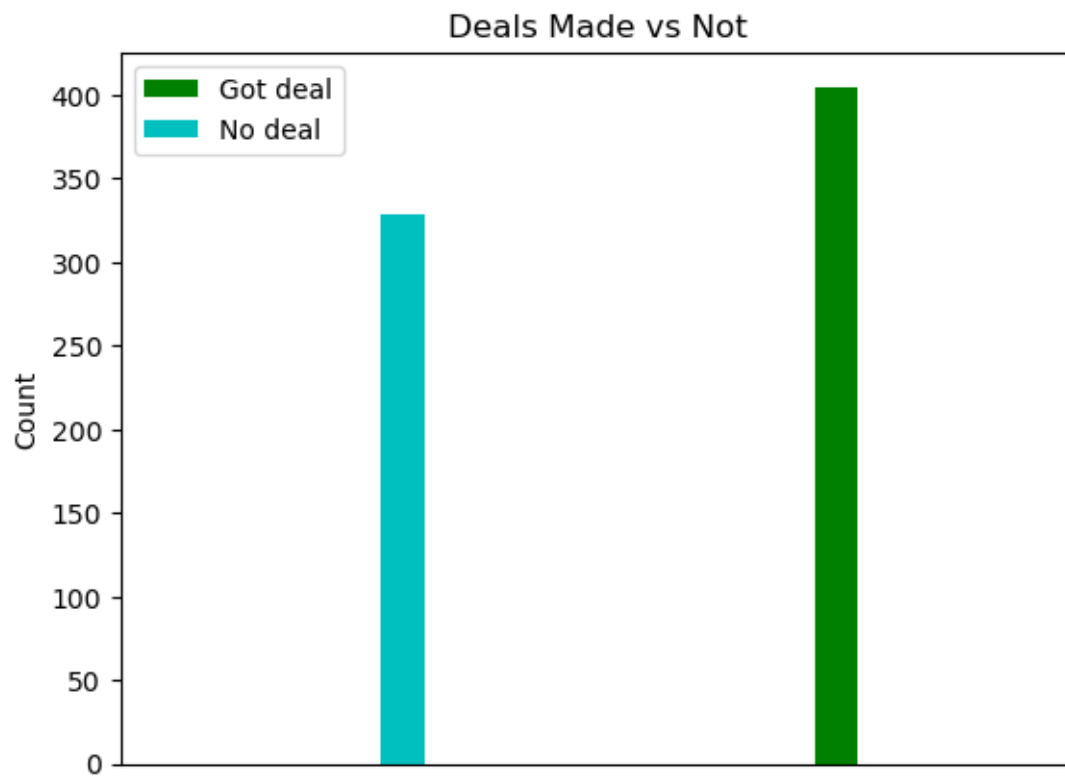
```
[11]: # Creating a dataset of only companies that did and did not get deals
yes_deal_df = dropping_nulls_df[dropping_nulls_df['Got Deal']==1]
```

```
no_deal_df = dropping_nulls_df[dropping_nulls_df['Got Deal']!=1]
```

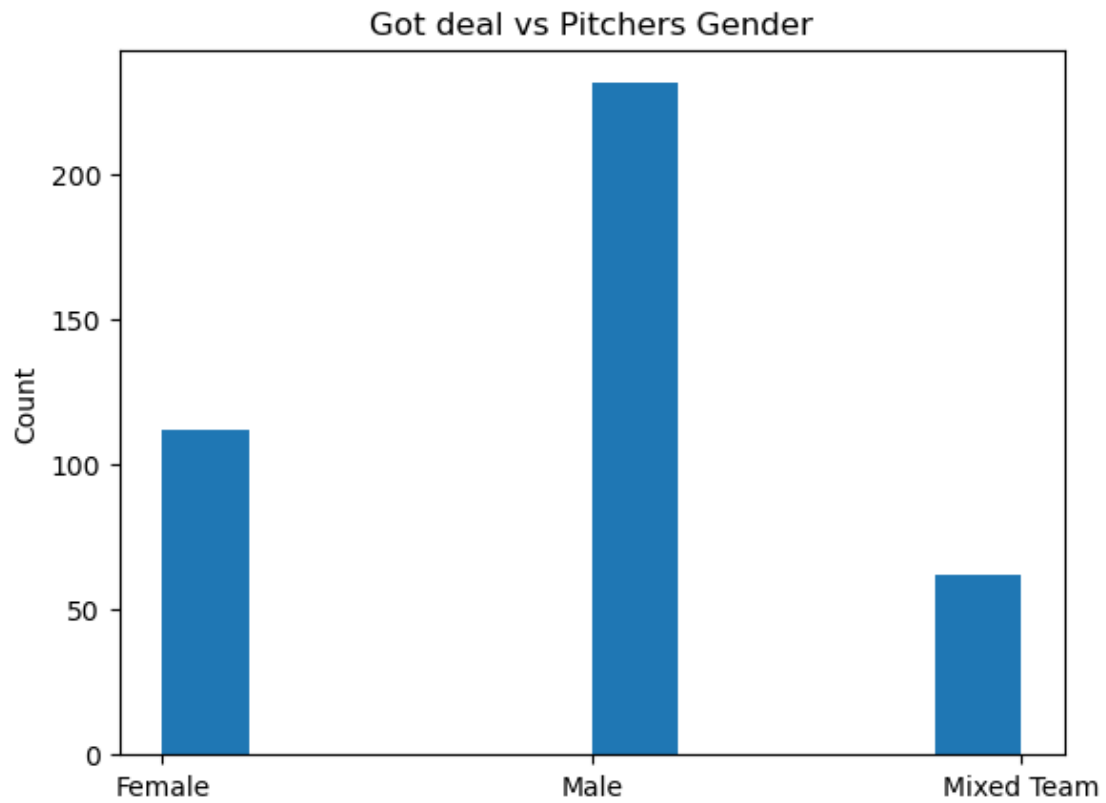
```
[12]: # Histogram of companies that did secure a deal and how much they asked for
fig, ax=plt.subplots()
# Using the below code to remove scientific notation from the x-axis
ax.ticklabel_format(style='plain')
ax.hist(yes_deal_df['Original Ask Amount'], bins=(50))
ax.set_title('Deals Made vs Original Asking Amount')
ax.set_xlabel('Ask Amount$')
ax.set_ylabel('Count')
plt.show()
```



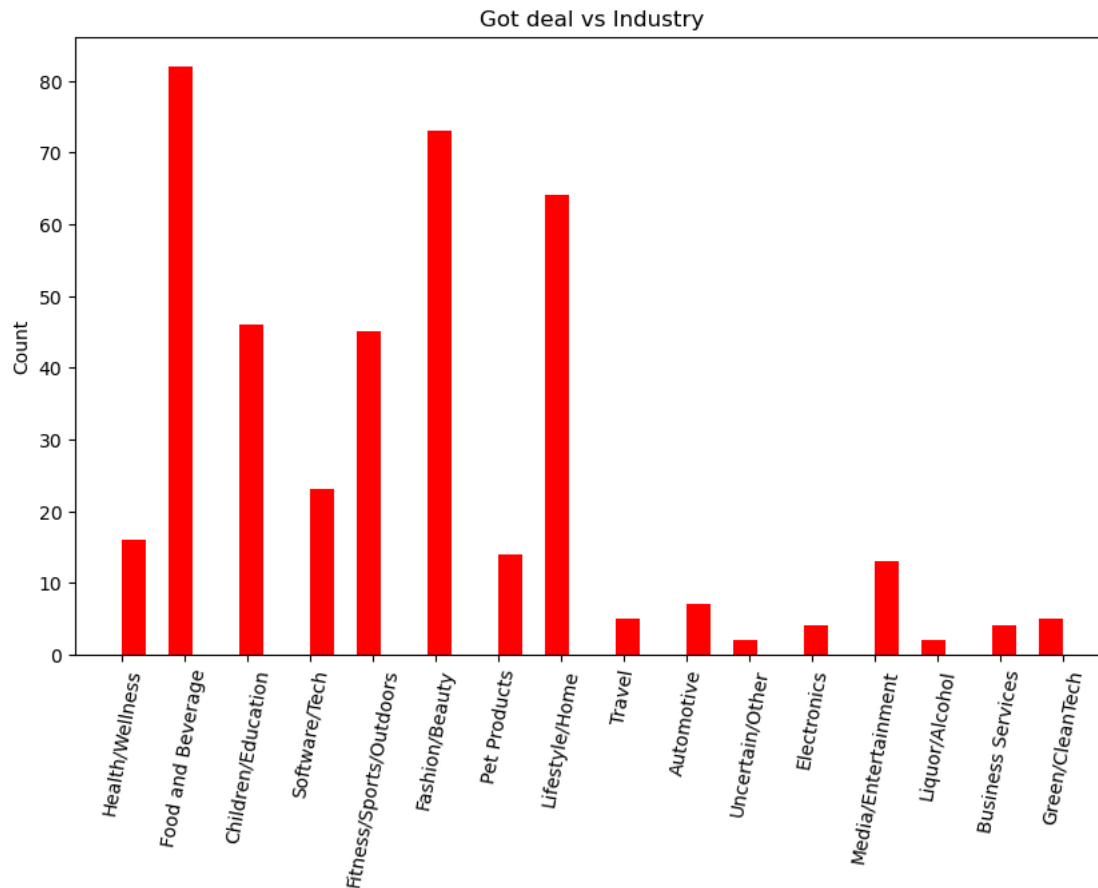
```
[13]: # Histogram of companies that did and did not secure deals
plt.hist(yes_deal_df['Got Deal'], color='g', label='Got deal')
plt.hist(no_deal_df['Got Deal'], color='c', label='No deal')
plt.legend()
plt.title('Deals Made vs Not')
plt.xticks([])
plt.ylabel('Count')
plt.show()
```



```
[14]: # Histogram of companies that secured deals by gender
plt.hist(yes_deal_df['Pitchers Gender'])
plt.title('Got deal vs Pitchers Gender')
plt.ylabel('Count')
plt.show()
```



```
[15]: # Histogram of deals made by industry
plt.figure(figsize=(10,6))
plt.hist(yes_deal_df['Industry'], bins=40, color='r')
plt.xticks(rotation=80)
plt.title('Got deal vs Industry')
plt.ylabel('Count')
plt.show()
```

3 Creating Models

3.0.1 Logistic regression

```
[16]: # Loading in the necessary packages
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, ConfusionMatrixDisplay
```

```
[17]: # Creating a feature and target dataframe
target_df = dropping_nulls_df['Got Deal']
feature_df = dropping_nulls_df.drop('Got Deal', axis=1)

# Creating a dummy variable dataframe to handle non-numeric features
dummy_df = pd.get_dummies(feature_df)
```

```
print(target_df.shape)
print(dummy_df.shape)
```

```
(734,)
(734, 72)
```

```
[18]: # Splitting the data
df_train, df_test, target_train, target_test = train_test_split(dummy_df,
    ↪target_df, test_size=30, random_state=200)
print(df_train.shape, df_test.shape, target_train.shape, target_test.shape)

# Scaling the data to cover the wide range of feature values
scaler = StandardScaler()
df_train_scaled = scaler.fit_transform(df_train)
df_test_scaled = scaler.transform(df_test)
```

```
(704, 72) (30, 72) (704,) (30,)
```

```
[19]: # Creating an instance of the regressor
log_reg = LogisticRegression(random_state=20).fit(df_train_scaled, target_train)


# Making a prediction of the y
predicted_test = log_reg.predict(df_test_scaled)
```

```
[33]: # Calculating metrics
cm = confusion_matrix(predicted_test, target_test)
accuracy = round(accuracy_score(predicted_test, target_test), 2)*100
precision = round(precision_score(predicted_test, target_test), 2)*100
recall = round(recall_score(predicted_test, target_test), 2)*100

print(f'The accuracy is: {accuracy}%\nThe precision is: {precision}%\nThe recall_
    ↪is: {recall}%')
```

```
The accuracy is: 77.0%
The precision is: 89.0%
The recall is: 77.0%
```

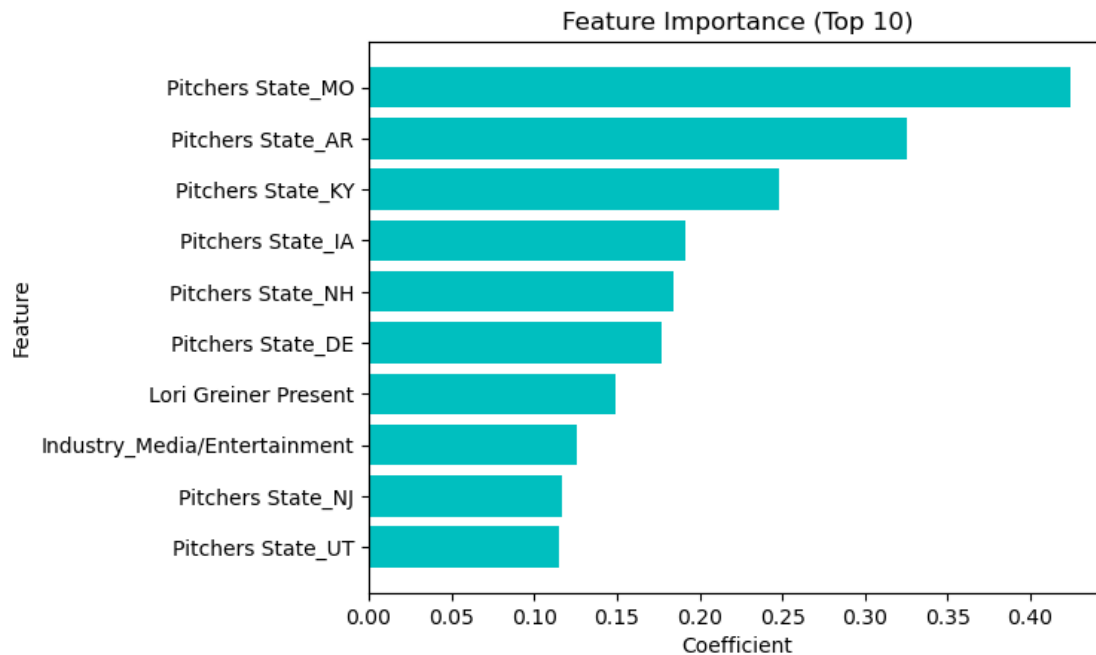
```
[34]: # Confusion matrix visualization
ConfusionMatrixDisplay(cm).plot(cmap='magma')
plt.show()
```



```
[22]: # Determining feature importance
      ## Setting the coefficient values and feature names to separate variables
      coefficients = log_reg.coef_[0]
      features = df_train.columns

      # Creating a dataframe for each feature and the respective f
      coef_df = pd.DataFrame(data=coefficients, index=features,
                             columns=['Coefficients']).sort_values(by=['Coefficients'],
                             ascending=False)
```

```
[23]: # Plotting the feature importance
      plt.barh(coef_df.index[:10], coef_df['Coefficients'].head(10), color='c')
      plt.gca().invert_yaxis()
      plt.title('Feature Importance (Top 10)')
      plt.xlabel('Coefficient')
      plt.ylabel('Feature')
      plt.show()
```



3.0.2 State related features appear to dominate so rerunning the model without state feature

```
[24]: # Creating a new dataset without pitcher state
n_s_df = explore_df[['Industry', 'Pitchers Gender', 'Original Ask Amount',
                    'Got Deal', 'Barbara Corcoran Present', 'Mark Cuban_
                    ↪Present', 'Lori Greiner Present',
                    'Robert Herjavec Present', 'Daymond John Present', 'Kevin_
                    ↪O Leary Present']]

# Dropping nulls
n_s_dropping_nulls_df = n_s_df.dropna()
n_s_dropping_nulls_df.shape
```

[24]: (895, 10)

```
[25]: # Creating a feature and target df
n_s_target_df = n_s_dropping_nulls_df['Got Deal']
n_s_feature_df = n_s_dropping_nulls_df.drop('Got Deal', axis=1)

# Creating a dummy variable dataframe
n_s_dummy_df = pd.get_dummies(n_s_feature_df)

print(n_s_target_df.shape)
print(n_s_dummy_df.shape)
```

```
(895,)
(895, 26)
```

```
[26]: # Splitting the data
n_s_train, n_s_test, n_s_target_train, n_s_target_test = \
    ↪train_test_split(n_s_dummy_df, n_s_target_df, test_size=30,
                    ↪random_state=200)

# Scaling the data
no_state_scaler = StandardScaler()
n_s_df_train_scaled = no_state_scaler.fit_transform(n_s_train)
n_s_df_test_scaled = no_state_scaler.transform(n_s_test)


# Creating an instance
n_s_log_reg = LogisticRegression(random_state=20).fit(n_s_df_train_scaled, \
    ↪n_s_target_train)
n_s_predicted_test = n_s_log_reg.predict(n_s_df_test_scaled)

[27]: # Calculating metrics
n_s_cm = confusion_matrix(n_s_predicted_test, n_s_target_test)
n_s_accuracy = round(accuracy_score(n_s_predicted_test, n_s_target_test), 2)*100
n_s_precision = round(precision_score(n_s_predicted_test, n_s_target_test), \
    ↪2)*100
n_s_recall = round(recall_score(n_s_predicted_test, n_s_target_test), 2)*100

print(f'The no state accuracy is: {n_s_accuracy}\n'
    ↪f'The no state precision is: {n_s_precision}\n'
    ↪f'The no state recall is: {n_s_recall}')
```

```
The no state accuracy is: 67.0
The no state precision is: 80.0
The no state recall is: 73.0
```

```
[35]: # Confusion matrix visualization
ConfusionMatrixDisplay(n_s_cm).plot( cmap='magma')
plt.show()
```



```
[29]: # Determining feature importance
      ## Setting the coefficient values and feature names to separate variables
      n_s_coefficients = n_s_log_reg.coef_[0]
      n_s_features = n_s_train.columns


      # Creating a dataframe for each feature and the respective f
      n_s_coef_df = pd.DataFrame(data=n_s_coefficients, index=n_s_features,
                                columns=['Coefficients']).
      ↪sort_values(by=['Coefficients'], ascending=False)

      # Plotting the feature importance
      fig, ax = plt.subplots(2, figsize=(6,6), sharex=True)
      ax[0].barh(n_s_coef_df.index[:10], n_s_coef_df['Coefficients'].head(10),
      ↪label=('Model excluding State'), color='r')
      ax[0].invert_yaxis()
      ax[0].legend()
      ax[0].set_title('Feature Importance (Top 10)')
      ax[0].set_ylabel('Feature')
```

```

ax[1].barh(coef_df.index[:10], coef_df['Coefficients'].head(10), label=('Model_
    including State'), color='g')
ax[1].invert_yaxis()
ax[1].set_xlabel('Coefficient')
ax[1].set_ylabel('Feature')
ax[1].legend()
plt.show()

```



4 Metric Recap

```

[36]: print(f'Confusion Matrix\n{cm}\n')
print(f'The accuracy is: {accuracy}%\nThe precision is: {precision}%\nThe recall_
    is: {recall}%\n', '-'*30, '\n', '-'*30)

print(f'No State Confusion Matrix\n{n_s_cm}\n')
print(f'The no state accuracy is: {n_s_accuracy}%\nThe no state precision is:_
    {n_s_precision}%\nThe no state recall is: {n_s_recall}%')

```

Confusion Matrix

```

[[ 6  2]
 [ 5 17]]

```

```
The accuracy is: 77.0%
The precision is: 89.0%
The recall is: 77.0%
```

```
-----
-----
No State Confusion Matrix
[[ 4  4]
 [ 6 16]]
```

```
The no state accuracy is: 67.0%
The no state precision is: 80.0%
The no state recall is: 73.0%
```

5 Conclusion

The original logistic regression model produced respectable metrics accross the board thus making the model useful for future use. Potential concerns are when it comes to the pitches states, high outcomes for respective states could be due to filming primarily being in the states thus limiting the ease of access for others. Additionally the final data size was 734 rows and 11 features. I would of liked to of had a minimum of 1000 rows; however, 734:11 is a decent ratio for modeling.